Starile unui proces:

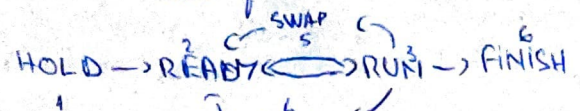

1. waits to be entered in system
2. loaded in mem. wait for CPU
3. runs the proces
4. wait for resources
5. when mem. is full the proc. is saved on HDD
6. the proc. is erased from mem.

Deadlocks - processes/threads enters an ∞ state bc. it wants a resource that is held by another

causes: 1. Mutual exclusion
2. Hold and wait - block a resource while waiting
3. Non-preemption - resources can't be stolen
4. Circular wait

Prevent: Pick an order for resources and always work / lock in that order

Detect: using the graph of processes
cycle = deadlock

Planificare:

1. First come first served
2. Shortest job first
   - we have to know the duration
   - not good for big clients
3. Priorities
   - starvation risk
4. Deadline scheduling
   - ed cere terminie să termine mai rpd.
5. Round Robin
   - eg. time periods
   - circular order
   - all in Ready / run

Loading methods:

V1: all prog. pages into RAM from the beg
+ execution is fast
- slow startup
- occupy data never used

V2: load 1st page, rest when needed
+ don't waste memory
- slow startup
- slow exec.

V3: Locality principle - a proc. is likely to need soon the pages next to the page just loaded
- prefetch neigh. pages

Threads vs Proc.

+ memory don't copy and stack for every T
+ quicker to create
+ The communication is better (shared mem.)
+ every T execute a function
- no security between T
- if one T blocks, all T block

Tehnici de alocare:

REAL
1. Mono-tasking - one task at a time
2. Multi-tasking:
   Fixed partitions:
   a) absolute
   + more proc., but limited nr.
   - limited dimension of a proc.
   - compiles the addr. of each pos.
   b) relocatable
   + no longer fixed RAM mem.
   - addresses are stored as an offset
   - more calcul. for each instr.
   Phy. addr = offset + part-start
   - less nr. of proc.
   Variable partitions:
   + no longer a predefined nr of part.
   + we can run any nr. of proc.
   - memory fragmentation

VIRTUAL
Paged
+ mem. frag solved
- use a page for a very small proc.
- more difficult calculus for addr.
(pag, offset)

Segmented
+ more mem than address size allowed
+ protect mem. access
+ reusable code
- mem. frag.

Paged-Segmented
- segments împărțite în mai multe pagini
(segm, page, offset)
- procesul are o tabelă de segm.
- fiec. tabelă de segm. are o tab. de pag.

How to choose pages for SWAP?

- NRU - page has 2 bits r/w
(0,0)(0,1)(1,0)(1,1) - se șterge mai întâi clasa 0, dacă nu există, dr1...
- FIFO - oldest page
- LRU - NFU matrix - delete the page with least 1 on line

Zombie = proc that has completed exec. but still has an entry in the proc. table
Use wait to avoid or before creating
signal (SIGCHLD, SIG_IGN)
Orphan = still exec. parent died

UNIX DISK

Block 0 - boot - small program
Block 1 - superblock - disk details
Block 2 - blocks for inodes
Block n
Block n+1 - blocks for data
Block max
Block - smallest size of data transferred

File types: regular, directory, hard/soft links, sockets, FIFO, peripheric coract./block
Symbolic links: ln -s target
- can reference on another disk
- delete file => points to a file that doesn't exist
Hard link: ln target
- new inode for the same file
- have to be root
- file is deleted only when there are no more links to it
=> we still have access to the data through our hard link.

How to handle malloc?
2 linked lists < free
occupied

- First-fit + fast
(first free space) - fragmentation
- Best-fit + recunoaște zonele mari de memorie
(best space) - slower
- creates small spaces of free mem => fragm.
- Worst-fit + leaves large free spaces
(biggest space) - slower
- Buddy-fit + speed
allocate a chunk of the smallest pow of 2, greater than the req. size; keeps list of free chunks by pow of 2 $2^m = 2^k + 2^k + 2^{k+1} + \dots + 2^{M-1}$

inode - entries point to the list of blocks of a file
1-10 primele blocuri de câte 512 bytes
11 - indirectoare simplă (conține adresele a 128 de blocuri a câte 512 bytes)
12 - indirectoare dublă
13 - indirectoare Triplă

$(5 + \frac{B}{A} + \frac{B}{A^2} + \frac{B}{A^3}) \cdot B$

B - size of block
A - size of address

AWK: NR - current line
NF - nr. of fields
$0 - entire input line
$1, $2, $3 fields in current line
awk -F: '{print $1}' file

cut -d',' -f1

Processes: ps -ef | awk '{print $1}' | sort | uniq -c | sort -n

wc: -l lines -w words -c characters

grep -E "." file
-c nr lines
-n print nr lines in file
-o only part that matches

sort -n reverse
-u eliminate duplicates

Memory hierarchies:
- registers
- cache - temporary memory
L1, L2, L3
- RAM
- SSD/HDD

L1: smaller, closer to the comp. unit faster
L2: slower, bigger
L3: typically, all the chips

where should we place the page k:
- Direct cache org.:
page k -> k % N
+ fast
- cache collisions (cache Trashing) => slower
- Set organising
first free pos., we have to iterate => slow
- Set associative
- organized in groups of pages
- calculated as page_nr % cache_set
+ flexibility
+ avoid collisions
- doesn't use all available cache lines

Read la FIFO: așteaptă până nu mai există procese deschise pt scriere sau până apar date; întoarce câți bytes a citit sau 0 dacă 3 procese vor să scrie

PIPE: + easy to work with
! Returns nr bytes written
int p[2]; pipe(p);
r/w(p[0/1], &val, sizeof(i))

FIFO: poate fi folosit de oricine dacă se știe calea și are permisiuni
mkfifo("fifo2k", 0666)
int fd2k;
fd2k = open("fifo2k", O_RDONLY)
w/r(fd2k, &val, sizeof(i))
No proc. to read/write:
wait until a proc. appears

SED
a) search and substitute
sed -E "s/([1-3])([1-3])/\2\1/gi" fis
g - global
i - case insensitive
b) translitterate
sed -E "y/aeiou/AEIOU" fis
c) delete
sed -E "/.../d" fis

Mutex: pthread_mutex_t mtx;
pthread_mutex_lock(&mtx);
pthread_mutex_unlock(&mtx);
pthread_t t[3];
pthread_create(&t[3], NULL, func, NULL);
pthread_join(t[3], NULL);
pthread_mutex_destroy(&mtx);

int dup(int vechi)
int dup2(int vechi, int nou)
dup2(p[2], 0);

Semaphore: sem_t name;
sem_init(&name, 0, val);
sem_wait(&name);
sem_post(&name);
sem_destroy(&name);
Binary sem.: has values 0/1, it works like a mutex and it is used to sync concurrent processes